

Naive Bayes

Jason Shipp

July 9, 2018

Concept

In an ideal world the value of any given feature would be totally independent from the value of another feature for a given data point. If this were the case we could use Bayes Law/Rule to compute the missing / predicted value with certainty after working out the subsequent relationship between the features. As it stands however this is not the case, and frequently there are complex underlying relationships in our data that are impossible to model on computers. It can, however, be useful to pretend that these relationships are independent and apply Bayes Law anyway. This is where Naive Bayes comes in!

Multinomial Naive Bayes

Math Involved

Bayes Rule

Let $x = (x_1, x_2, \dots, x_n)$ $p(C_k|x) = \frac{p(C_k)p(x|C_k)}{p(x)}$ Read $p(y|x)$ as the probability of y given that x has occurred (1)

Naive Bayes Assumption

$$C_k = \operatorname{argmax}_p(C_k) \prod_{i=1}^n p(x_i|C_k) \quad (2)$$

Code

Scikit-Learn

```
1 import pandas as pd
2 import numpy as np
3 from sklearn.datasets import load_iris
4 from sklearn.model_selection import train_test_split
5 from sklearn.feature_extraction.text import CountVectorizer
6 from sklearn.naive_bayes import MultinomialNB, GaussianNB
7 from sklearn import metrics, preprocessing
8
9 #SCIKIT-LEARN PORTION
10 #THE FOLLOWING CODE IS ADAPTED FROM RITCHIE NG AND DOES NOT BELONG TO THE REPO OWNER
11 #Resource: https://www.ritchieng.com/machine-learning-multinomial-naive-bayes-vectorization/
12 #Read in the data into a pandas dataframe
13 dataFile = 'https://raw.githubusercontent.com/justmarkham/pycon-2016-tutorial/master/data/sms.tsv'
14 features = ['label', 'message']
15 sms = pd.read_table(dataFile, header=None, names=features)
16
17 #If you want to see it's size
18 #texts.shape
19
20 #If you want to see its top 5 or 10 or whatever
21 #texts.head()
22
23 #If you want to see the unique labels
24 #texts.label.unique()
25
26 #Lets convert the discrete "label" into a numerical value we can work with
27 #Some examples will show you doing this as a .map and then doing it manually but this
28 #Should work in the general sense
29 le = preprocessing.LabelEncoder()
30 sms['labelNum'] = le.fit_transform(sms['label'])
31
32
33 X = sms.message
34 y = sms.labelNum
35
36 #Split into training and test states, you can set a random_state to have repeatable
37 #occurrences
38 X_train, X_test, y_train, y_test = train_test_split(X, y)
39 print(y_test.head())
40 #Here we are going to use a "Count Vectorizer" to create a bag of words model
41 #since we are dealing with textual data
42 #This will be a "sparse matrix" and its best to let sklearn/pandas handle this as
43 #it knows how to store it in an effecient manner
44 vect = CountVectorizer()
45 X_train_dtm = vect.fit_transform(X_train)
46 X_test_dtm = vect.transform(X_test)
47
48 #Let's create a model
49 nb = MultinomialNB()
50 nb.fit(X_train_dtm, y_train)
51
52 #Here is where we make an actual prediction on our test data, easy right
53 y_pred = nb.predict(X_test_dtm)
54
55 #How'd we do?
56 print('Multivariate Naive Bayes')
57 print('Accuracy: ', metrics.accuracy_score(y_test, y_pred))
58 print('Precision: ', metrics.precision_score(y_test, y_pred))
59 print('F1 Score: ', metrics.f1_score(y_test, y_pred))
60
61 #Top left is True Negative predictions, top right is False Positives
62 #Bottom left is False Negatives, bottom right is True Positives
63 print(metrics.confusion_matrix(y_test, y_pred))
```

```

64
65 #If you want to test the probability a sentence will have a certain category you can
66 #Use this, however this really isn't a very good use of naive bayes
67 #y_pred_prob = nb.predict_proba(X_test_dtm)[: , 1]
68
69 #You can also use Gaussian NB if our matrix wasn't sparse.
70 '''
71 gnb = GaussianNB()
72 gnb.fit(X_train_dtm, y_train)
73 y_pred = nb.predict(X_test_dtm)
74 print('Gaussian Naive Bayes')
75 print('Accuracy: ', metrics.accuracy_score(y_test, y_pred))
76 print('Precision: ', metrics.precision_score(y_test, y_pred))
77 print('F1 Score: ', metrics.f1_score(y_test, y_pred))
78 '''
79
80 #TENSORFLOW Portion
81 #This is usually not the use for Tensorflow, but you can kind of force it if you want
82 class NaiveBayesClassifier:
83     dist = None
84
85     def fit(self, X, y):
86         ydist = np.unique(y)
87         count = np.array([
88             [x for x,t in zip(X,y) if t == c]
89             for c in ydist])
90
91         mean, var = tf.nn.moments(tf.constant(count), axes=[1])
92         self.dist = tf.distributions.Normal(loc=mean, scale=tf.sqrt(var))
93
94     def predict(self, X):
95         assert self.dist is not None
96         nb_classes, nb_features = map(int, self.dist.scale.shape)
97         prob = tf.reduce_sum(
98             self.dist.log_prob(
99                 tf.reshape(
100                     tf.tile(X, [1, nb_classes]), [-1, nb_classes, nb_features])),
101                 axis=2)
102             priors = np.log(np.array([1.0/nb_classes] * nb_classes))
103             joint_likelihood = tf.add(priors, cond_probs)
104         #TODO
105         norm_factor = tf.reduce_logsumexp(joint_likelihood, axis=1,
106             keep_dims=True)
107         log_prob = joint_likelihood - norm_factor
108         return tf.exp(log_prob)
109
110 #TODO Use classifier, later

```
